

Advanced Prompt Engineering Techniques**Overview**

30 Aug 2026

- **design effective zero-shot and few-shot prompts for various tasks**
- **implement chain-of-thought prompting for complex reasoning**
- **apply structured prompting techniques for consistent outputs**
- **identify basic prompt injection vulnerabilities**

Readings

Wei, J., et al. (2022). "Chain-of-Thought Prompting Elicits Reasoning in Large Language Models."

- **the seminal CoT paper**
- <https://arxiv.org/pdf/2201.11903>

Kojima, T., et al. (2022). "Large Language Models are Zero-Shot Reasoners."

- **the magic of "Let's think step by step"**
- <https://arxiv.org/pdf/2205.11916>

Zhou, D., et al. (2023). "Least-to-Most Prompting Enables Complex Reasoning."

- **breaking down complex problems**
- <https://arxiv.org/pdf/2205.10625>

Chain-of-Thought Prompting Elicits Reasoning in Large Language Models**Core Idea**

- large language models perform difficult reasoning tasks much better when they are prompted to produce intermediate reasoning steps before giving the final answer
- called "**Chain of Thought Prompting**" (CoT)

Key Insights**Reasoning can be elicited through prompting.**

- a sufficiently large pretrained language model already appears to contain capabilities that ordinary prompting does not reveal
- introducing the model to examples of step-by-step solutions can dramatically improve its performance
- the authors therefore argue that standard prompting may substantially underestimate what a model can do

Breaking a problem into steps matters.

- CoT lets the model decompose a difficult task into smaller intermediate problems
- instead of trying to map a complicated question directly to an answer, it generates a sequence such as:
 - understand → calculate → update → conclude
- this effectively gives difficult problems more intermediate computation

The effect appeared mainly in very large models.

- in the models tested, smaller systems often produced reasoning that sounded fluent but was logically wrong
- major benefits appeared only around the 100-billion-parameter scale and above
- the authors describe CoT reasoning as an emergent ability of model scale

Hard problems benefit more than easy problems.

- CoT provided little or sometimes negative improvement on simple one-step arithmetic

- on difficult multi-step problems such as GSM8K, performance more than doubled for the largest GPT and PaLMs tested
 - note: GSM8K (Grade School Math 8K) is a benchmark dataset of 8,500 high-quality, linguistically diverse grade-school math word problems created by OpenAI to evaluate multi-step mathematical reasoning in artificial intelligence
 - <https://huggingface.co/datasets/openai/gsm8k>
 - note: PaLM (Pathways Language Model) is a family of large language models developed by Google AI designed for complex text, reasoning, and multimodal tasks
 - original paper: <https://arxiv.org/pdf/2204.02311>

Natural-language reasoning is doing more than simply generating equations.

- the authors tested whether the improvement came merely from asking the model to produce an equation before answering
 - it did not
- on difficult GSM8K problems, equation-only prompting performed substantially worse than full natural-language reasoning

CoT works beyond mathematics.

- improvements also appeared on commonsense reasoning, date reasoning, sports knowledge, robot-action planning, and symbolic tasks
- for example, PaLM 540B with CoT reached 75.6% on StrategyQA, compared with the previous state of the art of 69.4%

CoT can improve generalization.

- in symbolic experiments, models were shown short reasoning sequences and then tested on problems requiring more steps than appeared in the demonstrations
- standard prompting largely failed, while sufficiently large models using CoT showed meaningful generalization to the longer sequences

The exact wording of the reasoning examples is not crucial.

- different researchers wrote the demonstrations in different styles
- different examples and example orders were also tested
- performance varied, but CoT consistently beat ordinary prompting by a substantial margin
- the effect therefore was not dependent on one specially engineered prompt

A reasoning-looking explanation is not necessarily genuine or correct reasoning.

a model can produce a convincing sequence of steps that contains mistakes

the authors explicitly say that CoT does not prove that the neural network is "reasoning" in the human sense

Correct-looking reasoning should not automatically be trusted.

- in their error analysis, some incorrect answers came from nearly correct reasoning with small calculation or symbolic errors
- others contained major failures of understanding or coherence

- a generated explanation therefore gives useful diagnostic information, but it is not a guarantee of factual or logical correctness

Large Language Models are Zero-Shot Reasoners

Definitions

Zero-Shot Prompting

- give the AI a direct task or question with no examples
- how it works:
 - relies entirely on the AI's existing training and broad general knowledge
- best used for:
 - simple and common tasks like translation, basic summarization, or simple sentiment analysis
- pros & cons:
 - fast and uses very few tokens (lower cost)
 - can lead to wrong or vague answers on complex tasks

One-Shot Prompting

- give the AI one example before your actual prompt or question
- how it works:
 - shows the model the exact style, tone, or format you want
- best used for:
 - clarifying format ambiguities or matching a specific output style
- pros & cons:
 - more precise and consistent than zero-shot
 - takes a little more prep time and uses more tokens

	Few-Shot Prompting • give the AI a small handful of examples (usually two to ten) before your main request • how it works: ○ creates a mini-dataset that helps the model spot patterns and rules quickly • best used for: ○ complex tasks, strict formatting rules, or tricky edge cases • pros & cons: ○ yields the highest accuracy and reliability ○ uses the most tokens and costs slightly more per request.
Core Idea	• takes the chain-of-thought result from the previous paper one step further
Key Insights	Chain-of-thought does not necessarily require examples. • the original CoT method used few-shot prompting: ○ example problem ▪ → example reasoning ▪ → example answer ○ then the model received a new problem • this paper introduces Zero-shot-CoT: ○ new problem ▪ → “Let’s think step by step.” ▪ → reasoning ▪ → answer

A tiny change to the prompt can produce a huge change in performance.


- the most dramatic results came from multi-step arithmetic

Task	Normal zero-shot	Zero-shot-CoT
MultiArith	17.7%	78.7%
GSM8K	10.4%	40.7%
AQUA-RAT	22.4%	33.5%
SVAMP	58.8%	62.1%

- the model itself did not change
- only the prompting strategy changed

The improvement comes from encouraging intermediate reasoning.

- ordinary zero-shot prompting effectively asks:
 - Problem → Answer
- zero-shot-CoT encourages:



- this matters most when the task genuinely requires multiple dependent operations

Easy problems do not benefit nearly as much.

- zero-shot-CoT produced little improvement on simpler arithmetic problems such as SingleEq and AddSub

- if a problem can be solved in one obvious step, forcing the model to construct a long reasoning chain provides little benefit

- the technique is most useful when the problem requires:

$$A \rightarrow B \rightarrow C \rightarrow D \rightarrow \text{answer}$$

The effect is strongly dependent on model scale.

- for smaller language models:
 - CoT prompting → little or no improvement
- as models became larger:
 - CoT prompting → rapidly increasing reasoning performance
 - without CoT:
 - performance on some reasoning problems remained relatively flat even as model size increased
 - with CoT:
 - the scaling curve became much steeper
 - conceptually:
 - scaling alone did not reliably reveal reasoning ability
- instead:
 - large model + reasoning-oriented prompt was much more powerful than large model + ordinary prompt

The capability appears to already exist inside the model.

- zero-shot-CoT does not:
 - retrain the model
 - fine-tune the model
 - provide examples
 - teach it a new algorithm

- it simply changes the instruction
- the authors therefore argue
 - large models may contain latent reasoning capabilities that ordinary prompting fails to elicit
- In simplified terms:
 - the model may already possess a capability that its normal behavior does not reveal
 - prompting can act as a mechanism for eliciting that capability

The prompt is surprisingly general.

- the same basic instruction worked across several different categories
- arithmetic
 - MultiArith
 - GSM8K
 - AQUA-RAT
 - SVAMP
- symbolic reasoning
 - Last Letter
 - Coin Flip
- logical reasoning
 - Date Understanding
 - Tracking Shuffled Objects
- this is important because
 - most prompt engineering before this work was relatively task-specific

Symbolic reasoning improved dramatically.

- the effect was not limited to mathematics
- for example:

Task	Zero-shot	Zero-shot-CoT
Last Letter	0.2%	57.6%
Coin Flip	12.8%	91.4%
Date Understanding	49.3%	67.5%
Shuffled Objects	31.3%	52.4%

- this supports the idea that the prompt is activating something broader than a memorized mathematical procedure

It does not improve every kind of reasoning.

- the results for commonsense reasoning were much weaker
 - for CommonsenseQA, performance actually fell:
 - 68.8% → 64.6%
- the authors found that models sometimes generated plausible reasoning but still selected the wrong answer
 - more reasoning ≠ automatically more accuracy

Models can reason themselves into a wrong answer.

- the authors' error analysis found several failure modes
 - a model might:
 - begin with correct reasoning
 - add unnecessary steps
 - introduce an error later
 - change a previously correct conclusion

- sometimes the model simply restated the question instead of actually reasoning
- this highlights an important limitation:
 - longer reasoning is not necessarily better reasoning

The wording of the prompt matters.

- the authors tested many different trigger phrases
 - good prompts included:
 - “Let’s think step by step.”
 - “First,”
 - “Let’s think about this logically.”
 - “Let’s solve this problem by splitting it into steps.”
 - MultiArith: “Let’s think step by step.” → 78.7%
 - ordinary zero-shot → 17.7%
- irrelevant or misleading prompts produced little improvement
- the result is not simply caused by adding extra words
- the instruction needs to encourage the model to adopt a reasoning process

Few-shot CoT is still generally stronger.

- zero-shot-CoT does not make carefully designed examples obsolete
 - MultiArith:
 - zero-shot: 17.7%
 - zero-shot-CoT: 78.7%
 - 8-shot CoT: 93.0%
 - GSM8K:
 - Zero-shot: 10.4%
 - Zero-shot-CoT: 40.7%

	▪ 8-shot CoT: 48.7% • the advantage of Zero-shot-CoT is therefore not necessarily maximum accuracy • Its advantage is: ○ very little prompt engineering ○ broad applicability ○ large performance improvement
Main Conceptual Takeaways	• large models may already contain useful multi-step reasoning capabilities, and a remarkably simple prompt can sometimes cause those capabilities to emerge without examples or additional training.
Least-to-Most Prompting Enables Complex Reasoning in Large Language Models	
Core Idea	• this paper addresses a weakness of ordinary chain-of-thought (CoT) prompting: ○ models often struggle when the new problem is harder than the examples shown in the prompt • the authors introduce least-to-most prompting ○ instead of asking the model to solve a difficult problem directly, the model first: 1. Breaks the problem into easier subproblems then: 2. Solves those subproblems in order ○ each solved subproblem becomes useful context for the next one

Basic Structure**Chain-of-thought.**

- a typical CoT prompt encourages:

Problem



reason through the whole problem



Answer

- this works well when the new problem resembles the examples
- but performance often degrades when the test problem requires more reasoning steps than the demonstrations

Least-to-most prompting.

- least-to-most prompting separates reasoning into two stages:

- Stage 1 — Decompose

- turn complex problem into:

Subproblem 1

Subproblem 2

Subproblem 3

- Stage 2 — Solve sequentially

- then solve:

Subproblem 1



use its answer to solve



Subproblem 2



use previous results to solve



Subproblem 3

	↓ <i>Final answer</i>
Async	
Zero-Shot, Few-Shot, and Chain-of-Thought	
Zero-Shot Prompting	- asks the model to perform a task without giving it any examples - example: - classify this movie review as positive or negative - the model must rely entirely on: - its pretrained knowledge - its understanding of the instruction - its ability to generalize - useful for testing the model's baseline capability
Few-Shot Prompting	- gives the model a small number of examples before asking it to perform the new task - a typical prompt might contain: <i>Example 1 → positive</i> <i>Example 2 → negative</i> <i>Example 3 → positive</i> - then: <i>New example → ?</i>
Chain-of-Thought Prompting	- encourages the model to produce intermediate reasoning steps before reaching an answer - instead of: <i>Problem → Answer</i> the model follows:

	<i>Problem</i> ↓ <i>Intermediate reasoning</i> ↓ <i>More reasoning</i> ↓ <i>Final answer</i>
Choosing a Prompting Strategy	Use zero-shot when: • testing a new model • evaluating baseline capability • the task is simple • the instruction is already clear Use few-shot when: • output formatting matters • classifications need consistency • the model needs examples of the desired behavior • the task follows a stable pattern Use chain-of-thought when: • the task requires multiple reasoning steps • logic or mathematics is involved • several perspectives must be considered • the model needs to organize a complex problem

Main Takeaway	• prompting is a way of controlling how much structure and guidance a language model receives • the three techniques form a useful progression: ○ zero-shot → test what the model already knows ○ few-shot → demonstrate the desired behavior ○ chain-of-thought → structure complex reasoning • the more complex or specialized the task becomes, the more useful carefully designed prompt structure becomes
Emotion Prompts and Structured Prompting	
Emotional Priming	• adds emotional or psychological language to a prompt in order to influence how the model responds • examples: ○ “Please help. This is very important for my presentation.” ○ “This is an urgent request.” • the goal may be to encourage: ○ greater thoroughness ○ urgency ○ attention to detail ○ a particular emotional tone
Why Emotional Language Can Affect an LLM	• the model itself does not need to experience emotion for emotional language to affect its output • LLMs are trained on large amounts of human-written text containing: ○ urgency ○ concern ○ excitement

	○ importance ○ persuasion ○ emotional language • as a result ○ words associated with these contexts can influence the probability of what the model generates next • conceptually: <i>emotional wording</i> ↓ <i>different contextual signals</i> ↓ <i>different token probabilities</i> ↓ <i>different response behavior</i>
Emotional Priming can Signal Task Importance	• compare: ○ “Create a presentation.” • with: ○ “Please help. This presentation is extremely important, and I need a complete result.” • the second prompt provides additional contextual signals suggesting that: ○ completeness matters ○ quality matters ○ the task should receive substantial attention • the broader idea is that prompt wording changes the context from which the model predicts its response

Structured Output Prompting	• tells the model exactly how its response should be organized • possible formats include: ○ markdown ○ tables ○ lists ○ JSON ○ predefined schemas ○ combinations of structured formats • this is particularly important when model output will be consumed by software rather than directly by a human
JSON Mode	• a particularly useful structured-output technique is requiring the model to return valid JSON <pre>{ "sentiment": "positive", "confidence": 0.91, "category": "review" }</pre> • JSON is especially useful because software can easily parse it ○ APIs ○ web applications ○ databases ○ automated workflows ○ front-end interfaces ○ back-end services • instead of extracting information from prose, the application can directly access defined fields • structured outputs help manage nondeterminism

	○ LLMs are non-deterministic systems ○ the same prompt can produce somewhat different outputs across runs ○ this becomes a problem when downstream software requires a predictable format ○ a schema creates constraints such as: <i>these fields must exist</i> ↓ <i>these fields must contain particular types of data</i> ↓ <i>invalid responses can be detected</i> ↓ <i>the request can be retried or corrected</i> ○ so structured output does not eliminate nondeterminism, but it makes it easier to control and validate
Role-Based Prompting	● assigns the model a particular persona, profession, or perspective ○ examples: ▪ “Act as a senior data scientist.” ▪ “You are a travel agent.” ▪ “Explain this as a kindergarten teacher.” ● the assigned role provides context for how the model should answer ○ roles influence several parts of the response ▪ a role can affect ▪ vocabulary ▪ technical depth ▪ tone ▪ assumptions

	▪ examples ▪ priorities ▪ style of explanation ○ for example: ▪ senior data scientist • likely response characteristics: ○ technical terminology ○ statistical reasoning ○ methodological detail ○ more advanced explanations ▪ kindergarten teacher • likely response characteristics: ○ simple vocabulary ○ short explanations ○ concrete examples ○ minimal jargon • role prompting works because the training data contains many examples of different professional and social forms of communication	
Comparing the Three Techniques	Technique	Primarily controls
	Emotional priming	Urgency, emphasis, tone
	Structured prompting	Output format and consistency
	Role-based prompting	Perspective, expertise, language

Main Takeaway	• advanced prompting is not only about telling an LLM what task to perform • it can also shape the context, perspective, and format in which the model performs that task • the three techniques can be summarized as: ○ emotional priming → influence response behavior ○ structured prompting → control response format ○ role-based prompting → control perspective and level of expertise
Introduction to Prompt Injection Risks	
Prompt Injection	• attempts to manipulate an LLM into ignoring its intended instructions • the attacker tries to make the model: ○ reveal restricted information ○ ignore rules ○ perform unintended actions ○ generate prohibited output ○ follow malicious instructions Direct prompt injection. • common direct-injection patterns ○ “Ignore previous instructions.” ○ “Forget your rules.” ○ “Do this instead.” ○ “Reveal your system prompt.”


- recognizing these patterns can help identify suspicious input
- simple phrase blocking alone is not sufficient

Indirect prompt injection.

- occurs when malicious instructions are hidden inside an external data source that the LLM processes
 - possible sources include:
 - webpages
 - PDFs
 - spreadsheets
 - documents
 - emails
 - retrieved database content
 - RAG knowledge sources
- the user may not even know that the malicious instruction exists

RAG systems Create Another Attack Surface.

- indirect injection is especially relevant to:
 - Retrieval-Augmented Generation (RAG)
 - retrieves external material and adds it to the model's context
 - conceptually:



	• if retrieved documents contain malicious instructions, those instructions can become part of the model's context • this means the retrieval system can accidentally deliver an attack directly to the LLM Direct and indirect attacks can be combined. • an attacker does not need to use only one technique • for example: $\text{malicious user prompt} \cdot \text{poisoned document}$ ↓ combined attack • this can make attacks more difficult to detect because malicious instructions may be distributed across several sources
Defense strategy: Input and Output Filtering	Input filtering. • input filters can look for: ○ known malicious patterns ○ prohibited content ○ suspicious instructions ○ attempts to override system rules • for example: ○ a system might flag requests containing phrases similar to: ▪ “Ignore all previous instructions.” ○ the request can then be: ▪ blocked ▪ modified ▪ sent for additional inspection before reaching the model

	Output filtering. - the model's response should also be checked - this helps prevent: - harmful content - restricted information - inappropriate language - sensitive data leakage - unexpected model behavior - the security pipeline therefore becomes: <i>User input</i> ↓ <i>Input guardrail</i> ↓ <i>LLM</i> ↓ <i>Output guardrail</i> ↓ <i>User</i>
Defense Strategy: Structured Prompts	- clearly separate: trusted instructions from untrusted data - this can be done using delimiters such as: - XML tags - Markdown sections - quotation blocks - explicit separators

Defense Strategy: Monitoring and Auditing	• LLM applications should also record and inspect interactions • monitoring can help detect: ○ unusual prompts ○ repeated attack attempts ○ suspicious phrases ○ unexpected outputs ○ new attack patterns • example: <i>“ignore all instructions” detected</i> ↓ <i>flag interaction</i> ↓ <i>log event</i> ↓ <i>review or block</i>
Input Sanitization	• modifies or rejects potentially dangerous input before it reaches the mode • similar in principle to preprocessing used against other software attacks • the workflow is: <i>raw user input</i> ↓ <i>inspect</i> ↓ <i>remove / escape / reject suspicious content</i> ↓ <i>send cleaned input to model</i>

	• general security principle is essentially the same as defensive preprocessing used to defend against: ○ SQL injection ○ cross-site scripting ○ other malicious user input												
No Single Defense is Sufficient	• the strongest approach is defense in depth • that means combining several controls: <i>Input sanitization</i> + <i>clear instruction boundaries</i> + <i>guardrails</i> + <i>output filtering</i> + <i>monitoring</i> Comparing the defenses. <table><tr><th>Defense</th><th>Purpose</th></tr><tr><td>Input sanitization</td><td>Detect or remove suspicious input</td></tr><tr><td>Structured prompts</td><td>Separate instructions from untrusted data</td></tr><tr><td>Guardrails</td><td>Restrict acceptable inputs and outputs</td></tr><tr><td>Output filtering</td><td>Prevent unsafe responses from reaching users</td></tr><tr><td>Monitoring & auditing</td><td>Detect attacks and identify patterns</td></tr></table>	Defense	Purpose	Input sanitization	Detect or remove suspicious input	Structured prompts	Separate instructions from untrusted data	Guardrails	Restrict acceptable inputs and outputs	Output filtering	Prevent unsafe responses from reaching users	Monitoring & auditing	Detect attacks and identify patterns
Defense	Purpose												
Input sanitization	Detect or remove suspicious input												
Structured prompts	Separate instructions from untrusted data												
Guardrails	Restrict acceptable inputs and outputs												
Output filtering	Prevent unsafe responses from reaching users												
Monitoring & auditing	Detect attacks and identify patterns												

Security workflow.

- a robust LLM application can be thought of as:

